

Build a working AI-powered intake workflow

Time limit: 90 minutes

Return window: 48 hours

Stack: your choice

THE GOAL

Build a working AI-powered workflow that turns messy intake into structured, reviewable output — and then explain it to someone non-technical.

The stack is yours to choose. Use n8n, Make, Zapier, Power Automate, Python, or a custom app. What matters is that it actually runs.

The setup

You are given three sample intakes (synthetic — no real client data):

1. A voicemail transcript from a family member, rambling and out of order.
2. A completed web-form submission with clean but incomplete fields.
3. A referral email from another attorney that is deliberately missing key information.

Today, a staff member reads each one, decides how urgent it is, and retypes the details into the case system.

The four mission goals

All four are required.

01 • Structured extraction

Produce a machine-readable JSON record per intake with: potential client name, contact information, facility or provider named, date of incident, injury type, referral source, and urgency flag.

02 • Attorney-ready summary

A five-sentence plain-English summary an attorney can read in twenty seconds and know whether to act.

03 • Traceability and honesty

Every extracted fact must be traceable to the source text. Anything not present in the source must return **NOT FOUND** — never a guess. Show us the rule that enforces this.

04 • Prove it fails safely

Run your workflow on intake #3 (the incomplete one) and show us exactly what it does when the information isn't there.

What to send us

DELIVERABLE	WHAT IT NEEDS TO INCLUDE
The build	Workflow export, repo link, or screenshots — plus the prompts you used.
A Loom (3 min max)	Run it live, show the failure case, and explain one design decision and one tradeoff you made.
One-page enablement doc	Written for a paralegal, not an engineer: what this does, when to trust it, when not to, and who to tell when it breaks.